

Automated Data Configuration for NoSQL Systems

Kevin Lee

Akshay Mittal

Sachin Ravi

1 High-level Design

1. Will receive SQL query workload
2. Finding Interesting Table-subsets with selection conditions imposed (based purely on table-subsets occurring in queries) ¹
3. Do merging to get more options (to get table-subsets that are not explicitly in any query)
4. For each of these table-subsets, we calculate score across all queries. Suppose Q is the set of queries we are considering and $m \in M$, the set of all table-subsets we are considering, is one possible table-subset,

$queryCost(m, Q) = \text{cost to answer query using table-subset } m$

$updateCost(m, Q) = \beta \cdot \text{number of queries that update any table in } m$

$storageCost(m) = \alpha \cdot \text{size of storing } m$

$cost(m, Q) = queryCost(m, Q) + updateCost(m, Q) + storageCost(m)$

Add each cost to priority queue

5. Get k lowest costs out of priority queue and keep the table-subsets corresponding to these k lowest costs (if having picked smaller subset, ignore larger subset if it is considered later)

1.1 Finding Interesting Table-Subsets

$Cost(T) = \text{cost of all queries where subset } T \text{ occurs}$

$$WeightedCost(T) = \sum_i cost(Q_i) \cdot \frac{\text{sum of sizes of tables in } T}{\text{sum of sizes of all tables referenced in } Q_i}$$

Lemma 1. For $T_1 \subseteq T_2$, $Cost(T_1) \geq Cost(T_2)$.

Lemma 2. $WeightedCost(T) \geq C \Rightarrow Cost(T) \geq C \Leftrightarrow Cost(T) \not\geq C \Rightarrow WeightedCost(T) \not\geq C$

¹Might be interesting to try without selection conditions imposed and no merging

Algorithm 1 Algorithm for finding potential table-subsets in query workload

INPUT: C (baseline), $maxSize$ (max size for a table-subset)

```
 $S_1 \leftarrow \{T \mid T \text{ is a table-subset of size 1 with } Cost(T) \geq C\}$   
 $i \leftarrow 1$   
  
while  $i < maxSize$  and  $|S_i| > 0$  do  
   $i = i + 1$   
   $S_i = \{\}$   
   $G \leftarrow \{T \mid T \text{ is a table subset of size } i \text{ and } \exists s \in S_{i-1} \text{ such that } s \subset T\}$   
  for  $g \in G$  do  
    if  $Cost(g) \geq C$  then  $S_i = S_i \cup \{g\}$   
  end if  
end for  
end while  
 $S \leftarrow S_1 \cup S_2 \cup \dots \cup S_{maxSize}$   
 $R \leftarrow \{T \mid T \in S \text{ and } Weight(T) \geq C\}$   
return  $R$ 
```

When picking table-subsets, we want to pick subsets that are in expensive queries (because materializing these table-subsets will decrease a great cost). We pass in a threshold value of the minimum total cost we would like to satisfy when considering a certain table-subset. We start with single table-subsets and continually work our way up to higher table-subsets, while ensuring that we still satisfy the threshold.

1.2 Merging Table-Subsets

The problem with only using the previous procedure is that we only consider table-subsets that exist in some query in the workload. Oftentimes, there are table-subsets that don't occur in any query but are useful to materialize anyway.

[Insert Example]

Algorithm 2 Pairwise Merge

INPUT: T_1 , table-subset, T_2 table-subset

```
if  $T_1 \cap T_2 = \emptyset$  then return  $\emptyset$   
end if
```

```
 $newT \leftarrow$  union of  $T_1$  and  $T_2$ 's projection columns (only the ones on the intersection of  $T_1$  and  $T_2$ 's intersecting table-subsets), intersection of  $T_1$  and  $T_2$ 's selection conditions, and intersection of  $T_1$  and  $T_2$ 's table-subsets
```

```
return  $newT$ 
```

Algorithm 3 Algorithm for producing merged table-subsets

INPUT: T , set of table-subsets $R \leftarrow T$ **while** $|R| > 1$ **do** $M' = \emptyset$ **for** $R_1, R_2 \in R$ **do** $newR \leftarrow$ Pairwise Merge of R_1 and R_2 $M' \leftarrow M' \cup newR$ **end for****if** $M' = \emptyset$ **then**return $R - T$ **end if****for all** $m \in M'$ **do**Remove both parents of m from R **end for** $R \leftarrow R \cup M'$ **end while**return $R - T$

1.3 Pruning

1.4 Cost Model

The basic thing we need to write is $cost(q, T)$ where q is a normalized query and $T = \{T_1, \dots, T_k\}$ is a subset of tables with possible selection conditions imposed (the possible denormalized option).

 $cost(q, T)$